

Search:

1. BFS: The BFS search strategy explores the shallowest node in the search tree.
queue = collections.deque(); queue.popleft()
2. DFS: explores the deepest node in the search tree
queue.pop()
3. UCS: explores the cheapest node first
[(2, 'C'), (7, 'B'), (6, 'A')] heappop(frontier) -> (2, 'C')
4. Heuristic Function: $h(n)$
 - o Cost (estimate) of the cheapest path from the state at node n to a goal state
 - o If n is a goal node $h(n) = 0$
5. A*: orders by backward cost + forward cost
 $f(n) = g(n) + h(n)$:
g: add values on all paths to n , h : the value on n
6. A heuristic $h(n)$ is admissible (optimistic) if
 $0 \leq h(n) \leq h^*(n)$,
where $h^*(n)$ is the true cost to the nearest goal from n
7. A* is optimal if an admissible heuristic is used
8. Heuristic cost \leq actual cost for each arc:
 $h(a) - h(c) \leq \text{cost}(a \text{ to } c)$
9. Consequence of consistency: The f value along a path never decreases, $h(a) \leq \text{cost}(a \text{ to } c) + h(c)$
10. Constraint Satisfaction Problems (CSP):
 - (1) A set of variables, $X = \{X_1, \dots, X_n\}$
 - (2) A set of domains, $D = \{D_1, \dots, D_n\}$, where $D_i = \{v_1, \dots, v_k\}$ for each variable X_i
 - (3) A set of constraints C which specify allowable combinations of values

To solve a CSP we need to define a state space:

Each state is defined by an assignment of values to some or all variables $\{X_i = v_i, X_j = v_j, \dots\}$

11. Minimax: Always starts from MAX

MINIMAX(s) =

$$\begin{cases} \text{UTILITY}(s) & \text{if TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if PLAYER}(s) = \text{MIN} \end{cases}$$

12. Depth-Limit Search (DLS): Search only to a limited depth in the tree

13. α - β Pruning Algorithm: Min version

Consider Min's value at some node n

n will decrease (or stay constant) while the descendants of n are examined

Let m be the best value that Max can get at any choice point along the current path from the root: If n becomes

worse ($<$) than m , Max will avoid it, and Stop considering n 's other children

14. Expectimax Search: take weighted average of children

Markov Decision Processes:

1. Discounted rewards:

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

$$U([s_0, s_1, s_2, \dots]) = \sum_{t=0}^{\infty} \gamma^t R(s_t) \leq \sum_{t=0}^{\infty} \gamma^t R_{max} = R_{max} / (1 - \gamma)$$

2. Value Iteration:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

Reinforcing Learning:

1. Temporal Difference Learning:

Sample of $V(s)$: $\text{sample} = R(s, \pi(s), s') + \gamma V^\pi(s')$

Update to $V(s)$: $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + (\alpha)\text{sample}$

Same update: $V^\pi(s) \leftarrow V^\pi(s) + \alpha(\text{sample} - V^\pi(s))$

2. Q-Learning:

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma \max_{a'} Q_k(s', a') \right]$$

3. Now is evaluating Q-value, not value.

4. Exploration Functions: Takes a value estimate u and a visit count n , and returns an optimistic utility

$$Q(s, a) \leftarrow_{\alpha} R(s, a, s') + \gamma \max_{a'} Q(s', a')$$
$$Q(s, a) \leftarrow_{\alpha} R(s, a, s') + \gamma \max_{a'} f(Q(s', a'), N(s', a'))$$

5. Approximate Q-Learning:

transaction = (s, a, r, s')

$$\text{difference} = [r + \gamma \max_{a'} Q(s', a')] - Q(s, a)$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}] \quad \text{Exact Q's}$$

$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a) \quad \text{Approximate Q's}$$

Markov Models:

1. Marginal Distributions: $P(t) = \sum_s P(t, s)$

2. Conditional Distributions: $P(a|b) = P(a, b) / P(b)$

3. The product rule: $P(y) * P(x|y) = P(x, y)$

4. The chain rule:

$$P(x_1, x_2, x_3) = P(x_1)P(x_2|x_1)P(x_3|x_1, x_2)$$

5. Bayes' rule:

$$P(x|y) = \frac{P(y|x)P(x)}{P(y)}$$

6. Independence: $P(X, Y) = P(X) * P(Y)$

$$X \perp\!\!\!\perp Y$$


7. Conditional Independence:

X is conditionally independent of Y given Z iff any x, y, z

$$(1) P(x, y|z) = P(x|z) * P(y|z)$$

$$(2) P(x|z, y) = P(x|z)$$

8. Joint Distribution of a MM:



Joint distribution:

$$P(X_1, X_2, X_3, X_4) = P(X_1)P(X_2|X_1)P(X_3|X_2)P(X_4|X_3)$$

From the chain rule, every joint distribution over X_1, X_2, X_3, X_4 can be written as:

$$P(X_1, X_2, X_3, X_4) = P(X_1)P(X_2|X_1)P(X_3|X_1, X_2)P(X_4|X_1, X_2, X_3)$$

Assuming: $X_3 \perp\!\!\!\perp X_1 \mid X_2$ and $X_4 \perp\!\!\!\perp X_1, X_2 \mid X_3$

$$P(X_1, X_2, X_3, X_4) = P(X_1)P(X_2|X_1)P(X_3|X_2)P(X_4|X_3)$$


9. Mini-Forward Algorithm:

$$P(x_t) = \sum_{x_{t-1}} P(x_{t-1}, x_t)$$
$$= \sum_{x_{t-1}} P(x_t \mid x_{t-1})P(x_{t-1})$$

10. Stationary distribution:

$$P_\infty(X) = P_{\infty+1}(X) = \sum_x P(X|x)P_\infty(x)$$

Hidden Markov Models:



1. Joint distribution of an HMM:

$$P(X_1, E_1, X_2, E_2, X_3, E_3) = P(X_1)P(E_1|X_1)P(X_2|X_1)P(E_2|X_2)P(X_3|X_2)P(E_3|X_3)$$

2. Conditional Independence: Current observation

independent of all else given current state. eg:

$$E_1 \perp\!\!\!\perp X_2, E_2, X_3, E_3 \mid X_1$$

3. Filtering / Monitoring:

$$B_t(X) = P(X_t \mid e_1, \dots, e_t)$$

4. Passage of time:

(1) Assume we have current belief:

$$B(X_t) = P(X_t \mid e_{1:t})$$

(2) Then after one time step passes:

$$\sum_{x_t} P(X_{t+1} \mid x_t) P(x_t \mid e_{1:t})$$

$$P(X_{t+1} \mid e_{1:t}) =$$

$$B'(X_{t+1}) = \sum_{x_t} P(X' \mid x_t) B(x_t)$$

(3) Or:

(4) Beliefs get "pushed" through the transitions

5. Observations:

(1) Assume we have $B'(X_{t+1})$

(2) After evidence comes in:

$$P(X_{t+1} \mid e_{1:t+1}) = P(e_{t+1} \mid X_{t+1}) B'(X_{t+1})$$

$$(3) B(X_{t+1}) \propto_{X_{t+1}} P(e_{t+1} \mid X_{t+1}) B'(X_{t+1})$$

6. The forward algorithm:

We are given evidence at each time and wanna know

$$B_t(X) = P(X_t \mid e_{1:t})$$

$$= P(e_t \mid x_t) \sum_{x_{t-1}} P(x_t \mid x_{t-1}) P(x_{t-1}, e_{1:t-1})$$

Bayes Nets:

1. Arcs: encode conditional independence

2. Probabilities in BNs:

$$P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i \mid \text{parents}(X_i))$$


3. A path is active if each triple is active:

○ Causal chain $A \rightarrow B \rightarrow C$ where B is unobserved (either direction)

○ Common cause $A \leftarrow B \rightarrow C$ where B is unobserved

○ Common effect (aka v-structure) $A \rightarrow B \leftarrow C$ where B or one of its descendants is observed

4. Maximum Expected Utility (MEU): maximize utilities



5. D-Separation:

Check all (undirected!) paths between and X_i and X_j

• If one or more active, then independence not guaranteed

• Otherwise (i.e. if all paths are inactive), then independence is guaranteed

6. Preferences:

(1) $A > B$, prefer A to B, $U(A) > U(B)$

(2) $A \sim B$: indifferent, $U(A) = U(B)$

(3) $U(p_1, S_1; p_2, S_2; \dots; p_n, S_n) = \sum p_i \cdot U(S_i)$

NLP:

1. Word2Vec: Go through each position t in the text, which has a center word c and context words o

2. Use the similarity of the word vectors for c and o to calculate the probability of o given c , i.e., $P(o \mid c)$

3. maximize L :

$$L(\theta) = \prod_{t=1}^T \prod_{\substack{-m \leq j \leq m \\ j \neq 0}} P(w_{t+j} \mid w_t; \theta)$$

4. Or minimize J , which is average negative log likelihood:

$$J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(w_{t+j} \mid w_t; \theta)$$

5. Calculate $P(W_{t+j} \mid W_t; \theta)$:

(1) Use **two vectors** per word w :


○ v_w when w is a center word

○ u_w when w is a context ("outside") word

(2) Average both at the end to obtain a single vector representation per word

(3) Then for a center word c and a outside word o :

$$P(o \mid c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$



6. Parameters:

(1) Recall that all word vectors are squeezed into θ

(2) Assume we have d -dimensional vectors and V many words

(3) Every word has two vectors!

$$\theta = \begin{bmatrix} v_{\text{aardvark}} \\ v_{\text{abacus}} \\ \vdots \\ v_{\text{zebra}} \\ v_{\text{zyzzyva}} \\ u_{\text{aardvark}} \\ u_{\text{abacus}} \\ \vdots \\ u_{\text{zebra}} \\ u_{\text{zyzzyva}} \end{bmatrix} \in \mathbb{R}^{2dV}$$

7. Optimization: Gradient Descent

To train the model, we gradually adjust parameters to minimize J

○ We change the parameters by walking down the gradient (aka slope in 1D)

■ We are going to have to take the (partial) derivative! Steps:

(1) In gradient descent, we consider the rate of change of J with respect to θ , i.e. $\frac{\partial J}{\partial \theta}$

(2) Then we take steps proportional to the negative of the gradient

